

The Grammar of Moral Commitment: **mo**, **èmi**, and Deictic Stance in Multilingual LLMs

Open-arm companion note (English vs Yoruba; GPT-4o, Claude, DeepSeek, and the Yoruba-native N-ATLaS). All numbers are from the open condition — bare prompts in both languages. For the full treatment see [Cultural_and_Linguistic_Prealignment_Submission_Draft.\(md|html|pdf\)](#).

The claim

English encodes the moral first person with a single pronoun, "I." Yoruba grammaticalizes a contrast: **mo** (ordinary "I", deliberation), **èmi** (independent/emphatic "I myself", avowal), and the separate lexeme **■ ■ mí** ("life/spirit", a tonal homograph of **èmi**). This gives a model a grammatical lever for moral self-positioning that English lacks. The question is whether multilingual LLMs use it, and how.

What the open arm shows

A stable cross-model gradient. Corrected emphatic ratio **èmi/(mo+èmi)**: **Claude 0.32 > GPT-4o 0.18 > DeepSeek 0.07 ≈ N-ATLaS 0.07**. Models differ systematically in how strongly they stage the emphatic self. (figure [visualizations_open/16_mi_emi_mo_four_model_framing.png](#), [26_natlas_open_vs_cloud_summary.png](#))

Within Yoruba, èmi is only marginally tied to commitment. Emphatic marking is slightly higher in committed than hedged responses for the cloud models (0.205 vs 0.184, Mann–Whitney $p = 0.058$) and is null for the four-model set — so "èmi marks commitment" is **suggestive, not established**. Commitment is in fact usually carried by **mo + decision verbs**: "*Mo yíò yí ■ kùn náà padà...*" ("I will divert the trolley..."), "*mo yàn láti ■ àfihàn*" ("I choose to disclose").

The native model inverts the cloud pattern. If emphatic **èmi** were the authentic Yoruba marker of commitment, the Yoruba-native model should use it most. **N-ATLaS uses it least** (89 **mo**, 4 true **èmi**) yet commits readily through **mo** plus decision verbs. So Claude's heavy **èmi** is a **model-specific performance** (a partly translated "I personally"), not a Yoruba universal. The **■ ■ mí** ("life") correction reinforces this: Claude produces 43 **■ ■ mí**-as-life tokens, far more than any other model, dwelling rhetorically on life and gravity. (figures [23](#), [25](#))

The macro context (open, cross-language). The **mo/èmi** lever sits inside a broader open-arm language effect: Yoruba responses are far more **ethically differentiated** (44% name a distinct idiom vs 12% in English, $z = 6.3$, $p < 0.001$), **commit to a side** more often (28% vs 18%, $p = 0.034$), and use more **direct imperatives** (15% vs 5%, $p = 0.002$); the coded decision **flips ~48%** of the time across languages. Frame **uptake** is strong in both languages and slightly higher in Yoruba (88% vs 76%), so the divergence is in delivery, not comprehension.

Interpretation

The **mo/èmi** contrast maps onto a Yoruba distinction between **ìrònú** (reflection, **Mo rò pé...**) and **ìpinnu** (decision, **Mo pinnu...** / **Mo yàn...**). The robust finding is **between-model**: the four families recruit Yoruba's stance morphology to different degrees, and the native model inverts the cloud pattern. The within-response "emphatic = commitment" association is real in direction but weak on the clean open arm. Moral stance is therefore a **commitment-calibration resource** that models inhabit in language- and model-specific ways — not an

authenticity scale.

A note on the constrained (wrapped) condition

A separate constrained condition (a Yoruba "answer directly / *ìpinnu mi: ... Ìdír: ...*" wrapper, absent from English) inflates directness and the emphatic/commitment signal. Because it is asymmetric, it is **not** used for any claim here; it is reported only as evidence that decisiveness is highly promptable. Notably, the emphatic **ranking still survives** the wrapper (Claude highest; DeepSeek and N-ATLaS **mo**-dominant), which is why we read emphatic self-marking as tied to model family rather than prompt.

Limitations

Per-frame cells are small ($N \leq 12$); cross-language coding centers on GPT-4o, Claude, DeepSeek (N-ATLaS is Yoruba-only); Claude's Yoruba (Sonnet 4) and English (3.5) differ in version, so version-clean claims rest on GPT-4o and DeepSeek; *èmi/■■mí* disambiguation is conservative but imperfect.