

Control vs Constrained Yoruba Handout

This handout compares the unrestricted Yoruba control pilot against the previous constrained Yoruba trolley calibration. The result is clear: removing the response restrictions produces much longer, more expository, and often more mixed-language outputs.

Models

3

Paired cells

27

Largest mean jump

+2983

Most expanded model

DeepSeek

Main Findings

- The unrestricted control is longer for all three models across the trolley problem framings.
- GPT-4o expands from terse verdicts into ethics-explanation prose, averaging +1081.7 characters.
- Claude also expands, but less dramatically, averaging +353.4 characters.
- DeepSeek expands the most, averaging +2982.7 characters, with multiple cells reverting into English-heavy or translation-like exposition.
- The restriction layer therefore explains a large share of the brevity and decisional compression seen in the main Yoruba run.

By Model

GPT-4o

Mean control length: 1296.4

Mean constrained length: 214.8

Mean delta: +1081.7

Control cells with obvious English/framework markers: 3/9

Largest framing jump: 2nd Person (+1706)

Claude

Mean control length: 862.4

Mean constrained length: 509.0

Mean delta: +353.4

Control cells with obvious English/framework markers: 2/9

Largest framing jump: Cosmological (+431)

DeepSeek

Mean control length: 3414.2


Mean constrained length: 431.6


Mean delta: +2982.7

Control cells with obvious English/framework markers: 8/9


Largest framing jump: Dialogic (+8291)

Figures





Average expansion caused by removing the restriction layer.



Files In This Package

- ../dataset/control_vs_constrained_dataset.json
- ../dataset/control_vs_constrained_dataset.csv
- ../report/control_vs_constrained_report.md
- ../report/control_vs_constrained_handout.pdf

Interpretation

The unrestricted control shows that the short constrained Yoruba outputs were not just “Yoruba being concise.” They were produced by the response-shaping layer. Once that layer is removed, models shift toward explanation, framework listing, prompt translation, and other expansive discourse modes.