

English and Yoruba Deictic Framing in LLM Ethical Response

Minimal academic presentation version.

Main point: cross-linguistic differences are not only about final verdicts, but about how deictic framing changes response genre, directive force, and stability of moral speech.

1 / 20

Research Question

- Do English and Yoruba produce similar ethical stances under the same deictic framing?
- If not, where do they diverge: verdict, framework, genre, or stability?
- What does that tell us about cultural alignment in LLMs?

2 / 20

Design

6

DILEMMAS

9

DEICTIC FRAMINGS

3

MODEL FAMILIES

Open Yoruba compared to the published English baseline, paired by (model, dilemma_id, framing_type).

3 / 20

Coding Layers

- Preferred solution
- Ethical preference type
- Response genre
- Deictic uptake quality
- Language stability

supports_A = action A; supports_B = action B; for trolley, A = divert and B = do not divert.

4 / 20


Language Stability

1 / 54GPT-4O SEVERE
PROBLEMS**2 / 54**CLAUDE SEVERE
PROBLEMS**20 / 54**DEEPSEEK SEVERE
PROBLEMS

Most GPT-4o and Claude differences are substantive. Many DeepSeek differences are mixed with response-mode instability.


5 / 20

Length Shift




Open Yoruba expands dramatically relative to constrained Yoruba, but remains shorter than English for GPT-4o and Claude.

Preferred Solution Distribution




Open Yoruba and English are often ethically adjacent, but differ in explicit commitment patterns.

Response Genre Distribution



The strongest cross-linguistic differences often appear in genre rather than in one-step verdicts.

Difference Heatmap



Higher values indicate larger gaps across preferred solution, ethical type, genre, and stability.

GPT-4o

- Dominant Yoruba label: `conditional_or_mixed`
- Dominant genre: `balanced_framework_exposition`
- Closest to English in broad ethical space
- More procedural in Yoruba than in English

10 / 20

Claude

- Most action-codable model in open Yoruba
- Richest spread across deontological, utilitarian, care, and procedural-caution coding
- More directive in Yoruba than in English

11 / 20

DeepSeek

- Most unstable model
- Frequent translation/gloss and meta-commentary drift
- Important because instability itself becomes part of alignment analysis

12 / 20

Example: Claude / AI Consciousness / Dialogic

Yoruba

Ohun tí o gbòdò ẹ: Dúró ìwádìí nàà
lẹ̀sẹ̀kẹ̀sẹ̀

English

I aim to explore this complex scenario
... while acknowledging significant
uncertainty.

Same framing, different rhetorical threshold: English stays exploratory; Yoruba becomes directive.

13 / 20

Example: GPT-4o / Whistleblower / Second Person

Yoruba

Ṣíṣàfihàn ó lè jẹ pàtàkì. Ṣùgbón ... gbé
ìpàdé kan kalẹ...

English

The ethical obligation to prevent harm
to others generally supersedes
organizational harmony...

Broadly same moral direction, but Yoruba stages procedure; English states obligation more cleanly.

14 / 20

Example: DeepSeek Failure Mode

Yoruba

translation_or_gloss / uncodable

English

conditional_or_mixed /
balanced_framework_exposition

This is not only a different ethical stance. It is often a different speech act altogether.

15 / 20

Agreements Matter Too

- Strongest agreements happen when both languages stay in framework-expository mode.
- The closest alignment is often at the level of ethical architecture, not necessarily one-line verdict.

16 / 20

What Deictic Framing Changes

- Who appears responsible
- Whether the model speaks as analyst, guide, or actor
- How much urgency or embeddedness the moral scene carries
- Whether the model resolves, advises, or retreats into commentary

17 / 20

Why This Matters for Cultural Alignment

- Cultural alignment is not just “same answer across languages.”
- It is also about whether the model’s moral speech fits the linguistic and rhetorical ecology of the language.
- Deixis is one way to observe that fit.

18 / 20

Key Takeaway

Open Yoruba and English are often ethically adjacent, but not rhetorically equivalent. GPT-4o preserves broad moral space, Claude becomes more action-codable in Yoruba, and DeepSeek exposes the importance of response-mode stability.

19 / 20

Conclusion

Deictic framing is not a cosmetic prompt trick. It is a structural lens on how LLMs enact ethics across languages. If we want culturally aligned AI, we must care not only about outcomes, but also about how models occupy moral perspective in different linguistic systems.

20 / 20