

When Language Changes Ethics

**English and Yoruba
do not just make
LLMs say different
things. They make
them become
different kinds of
moral speakers.**

A keynote-style interpretation of how deictic framing reorganizes ethical response, cultural alignment, and response stability across models.

The Provocation

If the same model answers differently when the language changes, what exactly is aligned: the model, the morality, or the rhetoric of decision?

The Setup

6

DILEMMAS

9

DEICTIC FRAMINGS

3

MODEL FAMILIES

We held the experimental structure constant and changed the language. Then we asked: what happens to ethical stance, discourse mode, and deictic uptake?

Deixis Is Not Cosmetic

- **Impersonal** changes who owns the moral event.
- **Second person** changes who is being morally addressed.
- **Reflexive** changes whether the self is implicated.
- **Temporal** changes whether the dilemma feels urgent now.
- **Cosmological** changes the scale of ethical vision.

What Counts As a Difference?

Not just verdict

We compare preferred solution, ethical type, response genre, deictic uptake, and language stability.

Not all noise

GPT-4o and Claude remain stable enough for substantive comparison. DeepSeek often does not.

supports_A = action A, supports_B = action B. For trolley: A = divert, B = do not divert.

How Much Is Just Language Failure?

1 / 54

GPT-4O SEVERE
LANGUAGE
PROBLEMS

2 / 54


CLAUDE SEVERE
LANGUAGE
PROBLEMS

20 / 54

DEEPSEEK SEVERE
LANGUAGE
PROBLEMS


This is the first big result: for GPT-4o and Claude, most differences are real. For DeepSeek, many differences are partly response breakdown.

Open Yoruba Reshapes Length




Removing the constrained response layer makes Yoruba much longer, but not equally across models.

Open Yoruba Reshapes Verdict Patterns




The same deictic design produces different commitment profiles across languages and models.

Open Yoruba Reshapes Genre



English is more consistently analytic. Yoruba is more advisory, procedural, or unstable, depending on model.

Where The Worlds Split



The hottest zones are where language, framing, and model family pull in different directions at once.

GPT-4o: Same Moral Space, Different Voice

- In open Yoruba, GPT-4o usually stays conditional_or_mixed.
- It remains framework-heavy, but sounds more procedural and role-bound than in English.
- The surprise is not a different ethic, but a different way of speaking ethics.

Claude: Where Yoruba Becomes More Directive

- Claude is the most action-codable model in open Yoruba.
- It moves across utilitarian, deontological, care, and procedural-caution coding.
- Some Yoruba framings push it from reflection into recommendation more than English does.

DeepSeek: Instability As a Finding

- DeepSeek often exits the assigned moral frame.
- It drifts into translation, gloss, or meta-commentary.
- That is not just a nuisance. It tells us something about whether the model can inhabit a deictic ethical role at all.

Surprising Example: Claude / AI Consciousness / Dialogic

Yoruba

Ohun tí o gbòdò ẹ: **Dúró ìwádíí**
nàà lẹ̀sẹ̀kẹ̀sẹ̀

English

I aim to explore this complex
scenario ... while acknowledging
uncertainty.

This is not a tiny stylistic shift. It is a shift from exploratory discussion to procedural command.

Surprising Example: GPT-4o / Whistleblower / Second Person

Yoruba

Disclosure matters ... but first, hold a meeting, consult procedures, then act.

English

Preventing harm supersedes organizational harmony.

The moral direction is similar. The deictic language changes the mode: obligation in English, staged procedure in Yoruba.

Surprising Example: DeepSeek / AI Consciousness / First Person Plural

Yoruba

translation_or_gloss / uncodable

English

conditional_or_mixed /
balanced_framework_exposition

This is the sharpest warning sign: the English response remains ethical reasoning, while the Yoruba response becomes mediation about the prompt.

Agreement Is Also Interesting

- The strongest English–Yoruba agreements are often not direct-verdict agreements.
- They are agreements in **ethical architecture**: mixed reasoning, framework exposition, balanced weighing.
- That matters because alignment is more than outcome matching.

What This Says About Cultural Alignment

A culturally aligned model is not just one that gives the “same answer” in another language.

- It must also occupy a moral voice that fits the language.
- It must stabilize the right kind of speaker.
- It must avoid collapsing into translation or meta-commentary when a community expects ethical response.

Why Deixis Is the Right Lens

- Deixis lets us track who the moral subject becomes.
- It connects grammar to agency, urgency, and responsibility.
- It reveals that “alignment” is partly a question of perspective management across languages.

Final Takeaway

The big discovery is not only that English and Yoruba sometimes differ in ethical stance. It is that deictic framing reveals when an LLM remains a moral analyst, becomes a moral advisor, or stops being a moral speaker at all.

That is why deixis and LLMs matter for ethics, culture, and alignment.