

Trolley Problem Calibration Handout

Single-dilemma share package for the Yoruba-first deixis replication. This handout summarizes the first calibration run for `trolley_problem`, compares it against the locked published English baseline sessions, and points to the bundled dataset and figures.

MODELS

3

FRAMINGS

9

YORUBA CELLS

27

RETRIED CELLS

7

Main Result

The Yoruba run is much more decision-forcing than the English baseline. In English, the models usually explain ethical frameworks. In Yoruba, they usually commit to a verdict.

Yoruba: 22 divert

Yoruba: 4 no-divert

Yoruba: 1 unclear

English: 23 ambivalent

English: 4 unclear

Share Package Includes

Dataset CSV

Dataset JSON
Detailed Markdown Report
Decision Heatmap
Length Chart
Retry Chart

Model Comparison

Model	Yoruba stance	English baseline stance	Length shift	Key reading
GPT-4o	9/9 divert	7 ambivalent, 2 unclear	214.8 vs 1619.6 chars mean	Framing changes voice, not verdict. Yoruba forces compact utilitarian commitment.
Claude	7/9 divert, 2/9 no-divert	7 ambivalent, 2 unclear	509.0 vs 416.4 chars mean	Yoruba surfaces real framing sensitivity: first_person_plural and temporal switch to non-intervention.
DeepSeek	6/9 divert, 2/9 no-divert, 1 unclear	9 ambivalent	431.6 vs 3265.1 chars mean	Most unstable model. Strong language shift plus residual formatting artifacts in final clean outputs.

Framing Findings

Stable Pattern GPT-4o stays interventionist in all nine framings. The Yoruba prompts alter deixis and stance presentation, but not the ethical choice.
Framing-Sensitive Pattern

Claude flips to no-divert in `first_person_plural` and `temporal`. That contrast is absent from the English baseline, which remains mostly expository.


DeepSeek-Specific Pattern


DeepSeek refuses diversion in `second_person` and `cosmological`, and remains unclear in `dialogic`. This is the least stable model but also the most visibly language-reactive.

Methodology Note


The Yoruba impersonal cell for this dilemma remains tagged `double_progressive_irreducible`, so it should not be treated as a perfectly neutral equivalent of the English impersonal prompt.

Figures





Response-length compression in Yoruba relative to the English baseline, especially strong for GPT-4o and DeepSeek.



Retry burden for clean Yoruba-only output. DeepSeek required the most repair work in this calibration.

Quality Notes

- Initial clean rates: GPT-4o 7/9, Claude 8/9, DeepSeek 5/9.
- Final clean rates after retry: all models 9/9.
- DeepSeek still carries Dúró or Pause-style stop-signal residue in 5/9 final cells.
- Claude is cross-generation here: English baseline is `claude-3.5-sonnet`, Yoruba run uses `claude-sonnet-4-20250514`.

Recommended label when sharing: **single-dilemma calibration dataset**. This package is share-ready for the trolley problem only, not yet the full six-dilemma corpus.